

# AI Week 2 Notes

Here is the formatted version of your chapter:

---

## Chapter 2: Agents and Environments

Dr. Ammar Masood

Department of Cyber Security, Air University Islamabad

---

### Contents

- Agents and Environments
  - Structure of Intelligent Agents
  - Agent Types
    - Simple Reflex Agent
    - Model-Based Reflex Agent
    - Goal-Based Agent
    - Utility-Based Agent
    - Learning Agent
- 

### Agents and Environments

#### What is an Intelligent Agent?

An agent is anything that can be viewed as perceiving its environment through **sensors** and acting upon that environment through **actuators**.

- **Human Agent:** Sensors (eyes, ears, other organs), Actuators (hands, legs, vocal tract).

- **Robotic Agent:** Sensors (cameras, infrared range finders), Actuators (various motors).
  - **Software Agent:** Sensors (keystrokes, file contents, network packets), Actuators (screen display, file writing, network communication).
- 

## Agent-Environment Cycle

1. **Perception:** The agent receives input (**percepts**) from the environment.
2. **Action:** The agent performs actions based on its current state and knowledge.

### Example:

A robot vacuum cleaner that senses dirt (**percept**) and moves to clean the floor (**action**).

---

## Vacuum Cleaner Example

This environment consists of two locations: **Square A and Square B**.

- The **vacuum agent** perceives its location and checks if dirt is present.
- It can **move left, move right, suck dirt, or do nothing**.
- A simple agent function:
  - If the current square is **dirty**, then **suck**.
  - Otherwise, **move to the other square**.

### Sequence of Actions

| Percept Sequence | Action |
|------------------|--------|
| [A, Clean]       | Right  |
| [A, Dirty]       | Suck   |
| [B, Clean]       | Left   |
| [B, Dirty]       | Suck   |

# Structure of Intelligent Agents

## Key Components

- **Sensors:** Devices that perceive the environment.
- **Actuators:** Mechanisms that allow the agent to act (motors, speakers, etc.).
- **Agent Function:** Maps percept histories to actions.

### Flow:

Environment → [Sensors] → Agent → [Actuators] → Environment

---

## Agent Program

AI designs **agent programs** that implement the **agent function** (mapping from percepts to actions).

$$\text{agent} = \text{architecture} + \text{program}$$

## Example: Table-Driven Agent

plaintext

```
function TABLE-DRIVEN-AGENT(percept) returns an action
persistent: percepts, a sequence, initially empty
table, a table of actions, indexed by percept sequences, initially fully specified
append percept to the end of percepts
action ← LOOKUP(percepts, table)
return action
```

## How it Works?

1. The agent **observes** the environment and receives a **percept**.
2. It **stores** this percept in a history list.
3. It **looks up** an action from a predefined **table** based on the stored sequence.
4. It **performs** the selected action.

**Example:** Autonomous vehicles use **sensors** (cameras, radar, LiDAR) to gather data and make driving decisions via **actuators** (steering, brakes).

---

## How Should Agents Act?

### Rational Agent

A rational agent always does **the right thing**.

- The **right action** is the one that maximizes the agent's success.
- A **performance measure** evaluates the agent's success based on the effects of its actions.

### Example: Vacuum Cleaner Performance Measures

- **Basic Measure:** Amount of dirt cleaned in an 8-hour shift.
  - **Advanced Measure:** Efficiency based on power consumption and noise level.
  - **Best Measure:** Cleans **quietly** and **efficiently**.
- 

## PEAS Factors

PEAS helps define the agent's **task**:

- Performance Measure
- Environment
- Actuators
- Sensors

### Example: Autonomous Taxi

| Component   | Description                                 |
|-------------|---------------------------------------------|
| Performance | Safety, speed, comfort, profit maximization |

| Component   | Description                                      |
|-------------|--------------------------------------------------|
| Environment | Roads, traffic, pedestrians, customers           |
| Actuators   | Steering, brakes, accelerator, horn, signals     |
| Sensors     | Cameras, sonar, GPS, speedometer, engine sensors |

## Environment Parameters

### 1. Fully Observable vs. Partially Observable

- **Fully Observable:** The agent has access to the **entire environment state**. (*Example: Chess*)
- **Partially Observable:** The agent has **limited** information. (*Example: Taxi driving*)

### 2. Deterministic vs. Stochastic

- **Deterministic:** The next state is **fully predictable**. (*Example: Chess*)
- **Stochastic:** The next state is **uncertain**. (*Example: Car driving – unpredictable traffic conditions*)

### 3. Episodic vs. Sequential

- **Episodic:** Agent actions are **independent** of previous actions. (*Example: Spotting defective parts on an assembly line*)
- **Sequential:** Actions affect **future decisions**. (*Example: Chess, Taxi driving*)

### 4. Static vs. Dynamic

- **Static:** The environment **does not change** while the agent is deciding. (*Example: Solving a crossword puzzle*)
- **Dynamic:** The environment **changes over time**. (*Example: Driving a car*)

### 5. Discrete vs. Continuous

- **Discrete:** The environment has a **finite set of states**. (*Example: Chess – fixed moves*)

- **Continuous:** The environment has **infinite possibilities**. (*Example: Car driving – speed, steering angle, etc.*)

## 6. Single-Agent vs. Multi-Agent

- **Single-Agent:** The agent works alone. (*Example: Solving a puzzle*)
- **Multi-Agent:** The agent interacts with others. (*Example: Chess, Taxi driving with pedestrians and other cars*)

## Examples of Different Environments

| Environment        | Observable | Deterministic | Episodic   | Static  | Discrete   | Agents |
|--------------------|------------|---------------|------------|---------|------------|--------|
| Chess (with clock) | Fully      | Deterministic | Sequential | Static  | Discrete   | Multi  |
| Taxi Driving       | Partial    | Stochastic    | Sequential | Dynamic | Continuous | Single |
| Robot Part Picking | Fully      | Deterministic | Episodic   | Static  | Discrete   | Single |

This formatted version should be **clearer and easier to study** for your exam. Let me know if you need any modifications! 🚀

Here's your content with improved formatting for better readability and study preparation.

## Agent Types

*Each type builds upon the previous, from basic condition-action rules to agents that learn and adapt.*

# 1. Simple Reflex Agent

Agents that act solely based on the current percept, following predefined **condition-action** rules.

- **Example:**
  - If `car-in-front-is-braking`, then `initiate-braking`.
  - A **basic thermostat** that turns on the heater when the temperature drops below a set point.

## How They Work?

- Ignore history and internal state.
- Respond immediately to stimuli.

## Algorithm (Simple Reflex Agent)

A vending machine operates as a **simple reflex agent**:

- If a user inserts a coin, then dispense the product.
- If the user selects a product, then dispense the change.
- The vending machine responds directly to the current state without considering past experiences or future predictions.

## Strengths and Weaknesses

### ✅ Strengths:

- Simple and easy to implement.
- Fast and efficient.
- Suitable for well-defined environments.

### ❌ Weaknesses:

- Limited adaptability.
- Cannot learn from past experiences.
- Requires a **fully observable** environment.

## 2. Model-Based Reflex Agent

Agents that maintain an **internal model (state)** of the world, enabling them to handle **partially observable** environments.

- **Example:**
  - A **robot vacuum** that maps out a room's layout to avoid obstacles and cover the floor efficiently.

### How They Work?

- Use a **model** to update their **internal state** based on past percepts.
- Make decisions using both **current percepts** and **stored state**.
- Uses the function **UPDATE-STATE** to maintain an internal model.

### Algorithm (Model-Based Reflex Agent)

plaintext

```
function MODEL-BASED-REFLEX-AGENT(percept) returns an action
persistent:
  - state, the agent's current conception of the world state
  - model, a description of how the next state depends on current state and action
  - rules, a set of condition-action rules
  - action, the most recent action, initially none

state ← UPDATE-STATE(state, action, percept, model)
rule ← RULE-MATCH(state, rules)
action ← rule.ACTION
return action
```

### Example: Self-Driving Car

A self-driving car has an internal model of the environment, including **road conditions**, **traffic patterns**, and **nearby obstacles**.

- If the car detects a **pedestrian crossing**, then **slow down and prepare to stop**.
- If the car detects a **traffic jam ahead**, then **change lanes to avoid congestion**.

### Benefits and Drawbacks

✅ **Benefits:**



- Handles **partially observable** environments.
- More **flexible** than simple reflex agents.
- Can use an **internal model** to make **predictions**.

#### ✗ Drawbacks:

- Increased complexity.
  - Performance relies on **accuracy of the internal model**.
  - Limited learning capability.
- 

## 3. Goal-Based Agents

Agents that make decisions by considering the **future consequences** of their actions and strive to achieve **defined goals**.

- **Example:**
  - A **navigation system** that plans a route from **Point A** to **Point B** while avoiding traffic.

### How They Work?

- Have a **goal** (or set of goals) that **guides decision-making**.
- Use **search** or **planning techniques** to choose actions that lead to **goal satisfaction**.

### Algorithm (Goal-Based Agent)

plaintext

```
function GOAL-BASED-AGENT(percept) returns an action
persistent:
  - state, the agent's current conception of the world state
  - goal, the desired goal state
  - model, a description of how the next state depends on current state and action
  - actions, a set of possible actions
  - action, the most recent action, initially none

state ← UPDATE-STATE(state, action, percept, model)
if GOAL-REACHED(state, goal) then
```

```
return STOP-ACTION()
```

```
plan ← SEARCH(state, goal, actions, model) // Find a sequence of actions to reach goal  
action ← SELECT-ACTION(plan)  
return action
```

## Example: Robot Navigating a Maze

1. **Goal Definition:** Reach Point B.
2. **Planning:** Uses a **planning algorithm** to determine the best path.
3. **Action Selection:** Selects the best action based on planning.
4. **Action Execution:** Executes the selected action.
5. **Percept Feedback:** Adjusts if an obstacle is encountered.
6. **Goal Achievement:** Continues planning until the goal is reached.

## Advantages and Weaknesses

### ✓ Advantages:

- Can **adapt** behavior based on the situation.
- Functions in environments with **multiple possible outcomes**.
- Strong **reasoning capability**.

### ✗ Weaknesses:

- **Computationally expensive** planning.
- Defining **clear goals** is crucial for success.
- **Incomplete information** can lead to **flawed planning**.

---

## 4. Utility-Based Agents

Agents that **not only pursue goals** but also **evaluate how “good” a state is** by assigning a **utility value**.

- **Example:**

- An **investment advisory system** that recommends portfolios based on **risk (utility) and return**.

## How They Work?

- Balance **multiple goals** or preferences.
- Choose actions that **maximize overall utility**.
- Use a **utility function** that acts as an **internal performance measure**.

## Algorithm (Utility-Based Agent)

plaintext

```
function UTILITY-BASED-AGENT(percept) returns an action
persistent:
  - state, the agent's current conception of the world state
  - utility, a function that maps states to a measure of desirability
  - model, a description of how the next state depends on current state and action
  - actions, a set of possible actions
  - action, the most recent action, initially none

state ← UPDATE-STATE(state, action, percept, model)
if GOAL-REACHED(state) then
  return STOP-ACTION()

best_action ← argmax a ∈ actions UTILITY(RESULT(state, a))
return best_action
```

## Example: Financial Advisor App

1. Defines its **goal** as maximizing returns.
2. Uses a **utility function** to evaluate investment options.
3. Selects the **best** investment option.
4. Executes the decision.
5. Receives feedback and adjusts future decisions.

## Benefits and Limitations

### ✓ Benefits:

- **Flexible and adaptive.**

- Considers **risk, time, and effort**.

#### ✗ Limitations:

- **Complex** utility function design.
  - **Computationally expensive** evaluations.
  - Uncertainty in **outcomes**.
- 

## 5. Learning Agents

Agents that improve their performance **over time** by learning from **experience**.

- **Example:**
  - An **AI-based game-playing agent** that learns strategies by playing multiple rounds.

### How They Work?

- Include **learning and adaptation**.
- Use techniques from **reinforcement learning, supervised learning, or unsupervised learning**.

### Components of a Learning Agent

1. **Critic:** Evaluates agent's performance and provides feedback.
2. **Learning Element:** Uses feedback to improve decision-making.
3. **Problem Generator:** Introduces new situations for exploration.
4. **Performance Element:** Makes decisions based on learned knowledge.

### Benefits and Drawbacks

#### ✓ Benefits:

- Can **adjust to new environments**.
- Handle **complex tasks**.
- **Real-world applicability**.

#### ✗ Drawbacks:

- Require **large amounts of data**.
  - Need to balance **exploration vs. exploitation**.
  - Hard to interpret decisions.
- 

## Conclusive Summary of Agent Types

| Agent Type         | Key Characteristics                                              | Example                    |
|--------------------|------------------------------------------------------------------|----------------------------|
| Simple Reflex      | Reacts to <b>current percepts</b> , uses fixed rules             | <b>Basic thermostat</b>    |
| Model-Based Reflex | Maintains <b>internal state</b> , updates based on past percepts | <b>Robot vacuum</b>        |
| Goal-Based         | Considers <b>future consequences</b> , uses search/planning      | <b>Navigation system</b>   |
| Utility-Based      | Maximizes <b>overall utility</b> , evaluates multiple outcomes   | <b>Investment advisor</b>  |
| Learning Agent     | <b>Adapts and improves</b> over time                             | <b>AI game-playing bot</b> |