

# AI Week 5 Notes

Here's your formatted document with improved readability while keeping the original wording intact.

---

## Week 5: Local Search Strategies

Dr. Ammar Masood

Department of Cyber Security,

Air University Islamabad

CS340 - Introduction to AI

---

### Table of Contents

- Local search strategies
  - Hill climbing
  - Random restart hill climbing
  - Simulated annealing
  - Local beam search
  - Gradient descent
- 

### Problem Statement

In the search problems of **Chapter 3**, we aimed to find paths through the search space, such as from **Arad to Bucharest**. However, sometimes we are only interested in the **final state**, not the path taken.

## Example

- **8-Queens Problem:** Finding a valid final configuration of 8 queens.
- **Other Applications:** Integrated-circuit design, factory floor layout, job shop scheduling, automatic programming, etc.

## Solution: Local Search Algorithms

- Use a **single current state** and move to neighboring states.
  - **Advantages:**
    - Uses very little memory.
    - Often finds reasonable solutions in large/infinite state spaces.
    - Useful for pure **optimization problems** (finding the best state).
- 

## Local Search / Optimization

- **Goal:** Find the best state based on an **objective function**.
  - **Objective function:** Measures the “fitness” of a state.
  - **Problem:** Find the **optimal state** that maximizes/minimizes the objective function.
- 

## Local Search Strategies

- Local search algorithms operate by **searching from a start state to neighboring states**, without tracking the path.
- They are **not systematic** and may fail to explore parts of the search space.
- **When is local search used?**
  - When the **state-space is too large** to explore completely.
  - Starts with an **initial solution** and iteratively improves it.
  - Unlike **uninformed search**, it **does not maintain a search tree**.
  - Solves **optimization problems** where the goal is to find the **best state**.

## State-Space Landscape

- If elevation = **objective function**, the goal is to find the **highest peak** (global maximum) → Hill Climbing.
- If elevation = **cost**, the goal is to find the **lowest valley** (global minimum) → **Gradient Descent**.

## Advantages

1. Uses **very little memory**.
  2. Can often find **reasonable solutions** in **large/infinite state spaces**.
- 

## Hill Climbing

### Idea

- Always move towards a **state with a higher value** (maximization) or **lower cost** (minimization).

### Limitations

- Gets stuck in local optima.
- Susceptible to plateaus and ridges.

### Examples

- **Robotics Pathfinding**: A robot moves toward higher ground to avoid obstacles.
  - **Traveling Salesman Problem (TSP)**: Finding a shorter route step by step.
- 

## Hill Climbing Algorithm

1. Start anywhere.
2. Always choose the **best neighbor**.
3. If **no better neighbors**, quit.

4. **Random selection** among the best successors (if multiple exist).

## Working

- Keeps track of **one current state**.
- Moves to the **neighboring state with the highest value**.
- Terminates when reaching a "**peak**" (no neighbor has a higher value).
- Does not look beyond **immediate neighbors**.

## Example: 8-Queens Problem

- **Initial State:** 8 queens placed randomly on the board.
  - **Successor Function:** Moving one queen per column ( $8 \times 7 = 56$  successors).
  - **Heuristic Function:** Number of **pairs of attacking queens**.
- 

## Drawbacks of Hill Climbing

### Problems

- **Local Maxima:** Gets stuck in **suboptimal solutions**.
- **Plateaus:** Flat regions with no direction for movement.
- **Ridges:** Higher solutions exist, but cannot be reached directly.

### Solutions

- Introduce **randomness** to escape local optima.
- 

## Variants of Hill Climbing

### 1. Stochastic Hill Climbing

- Chooses **randomly** from uphill moves.
- Slower convergence but may **find better solutions**.

### 2. First-Choice Hill Climbing

- Randomly generates successors until one is **better than the current state**.
  - **Useful when many successors exist.**
- 

## Random Restart Hill Climbing

- **Idea:** "If at first, you don't succeed, try again."
- Conducts multiple hill-climbing searches from **randomly generated initial states**.
- **Complete with probability 1** (eventually finds a goal state).
- **Expected number of restarts** =  $1/p$  (where  $p$  is the probability of success).

### Example

- Climbing a **mountain range**:
    - Start climbing randomly → may reach **local maxima**.
    - Restart from a **new location** to find a **higher peak**.
- 

## Simulated Annealing

- **Inspired by metal cooling** → Gradually reducing randomness.
- **Accepts worse solutions** with some probability → Avoids local optima.
- **Uses a cooling schedule** (temperature  $T$  **gradually decreases**).

### Example Applications

- **Job Scheduling:** Finding the best sequence to reduce delays.
- **Neural Network Training:** Avoids poor convergence to bad local minima.

### Working

1. If move improves the state, always accept.
2. If move worsens the state, accept with **some probability**.
3. **Probability decreases as:**

- Move gets worse ( $\Delta E$  increases).
- Temperature (T) decreases.

## Cooling Schedule

- If T is reduced too fast, solution quality suffers.
  - $T(t) = a * T(t-1)$ , where  $0.8 \leq a \leq 0.99$ .
- 

## Local Beam Search

- Expands only the **k best nodes** at each level (unlike BFS/DFS).
- Tracks the **most promising paths**.
- **More efficient** than exhaustive search, but may **miss the global best solution**.

## Working

- Tracks **k states instead of one**.
- At each step:
  - Generates **all successors** of k states.
  - Selects the **k best successors**.
  - If a goal state is found, return it.

## Comparison with Random Restart

- **Random Restart**: Each search is independent.
  - **Local Beam Search**: Best successors influence other searches.
- 

## Genetic Algorithms

- Inspired by **evolutionary biology**.
- **Faster, randomized search** for optimal solutions.

## Operations

- **Crossover:** Combines two parents to create new children.
- **Mutation:** Introduces random changes.

## Working

1. **Initialize population** (random solutions).
  2. **Crossover & Mutation** on the fittest individuals.
  3. **Keep only the best solutions.**
- 

## Gradient Descent

- **Optimization algorithm** to minimize a **cost function**.
- Iteratively **adjusts parameters** to reduce error.

## Working

1. **Initialize parameters** (random values).
  2. **Compute Gradient:** Find direction of steepest increase.
  3. **Update Parameters:** Move **opposite to gradient**.
- 

This formatting improves readability and ensures all the original details are maintained. Let me know if you need further adjustments! 🚀