

Problem Solving Search

Here's your chapter with improved formatting for better readability:

Chapter 3: Problem Solving Agents

Dr. Ammar Masood

Department of Cyber Security, Air University Islamabad

Contents

- Problem Solving by Searching
 - Problem Formulation
 - Search Strategies
 - Uninformed Search Strategies
-

Problem Solving as Search

One way to address these issues is to view goal attainment as problem solving, and to see problem solving as a search through a state space.

For example, in chess, a *state* is a board configuration.

Important Concepts in Problem Solving

- The solution to any problem is a fixed sequence of actions.
- The process of looking for a sequence of actions that reaches the goal is called **search**.
- A **search algorithm** takes a problem as input and returns a solution in the form of an action sequence.

- Once a solution is found, the actions it recommends can be carried out. This is called the **execution phase**.
 - After formulating a goal and a problem to solve, the agent calls a **search procedure** to solve it.
-

Problem Solving by Searching

A problem can be formally defined by five components:

1. **Initial state** – The state the agent starts in.
 2. **Actions** – A description of the possible actions available to the agent.
 3. **Transition model** – Describes what each action does.
 4. **Goal test** – Determines whether a given state is a goal state.
 5. **Path cost function** – Assigns a numeric cost to each path.
-

A Simple Problem-Solving Agent

Algorithm:

sql

```
function SIMPLE-PROBLEM-SOLVING-AGENT(percept) returns an action
  persistent:
    seq, an action sequence, initially empty
    state, some description of the current world state
    goal, a goal, initially null
    problem, a problem formulation

  state ← UPDATE-STATE(state, percept)
  if seq is empty then
    goal ← FORMULATE-GOAL(state)
    problem ← FORMULATE-PROBLEM(state, goal)
    seq ← SEARCH(problem)
    if seq = failure then
```

```
        return a null action
    action ← FIRST(seq)
    seq ← REST(seq)
    return action
```

End Goal?

- A solution to a problem is an **action sequence** that leads from the initial state to a goal state.
 - **Solution quality** is measured by the **path cost function**, and an **optimal solution** has the lowest path cost among all solutions.
-

Example Problems

Toy Problems (Illustrate problem-solving methods)

- Concise, exact description
- Can be used to compare performance
- Examples:
 - 8-Puzzle
 - 8-Queens Problem
 - Cryptarithmic
 - Vacuum World
 - Missionaries and Cannibals
 - Simple Route Finding

Real-World Problems (Practical application)

- More difficult, often without a single, agreed-upon description.
- Examples:
 - Route Finding
 - Touring and Traveling Salesperson Problem (TSP)
 - VLSI Layout

- Robot Navigation
 - Assembly Sequencing
-

Problem Formulation

- **Toy Problem:** Designed for research and algorithm comparison.
- **Real-World Problem:** Solutions have practical applications but often lack a single precise description.

Example: Vacuum World

- **States:** Determined by agent location and dirt locations.
 - **Initial State:** Any state can be designated as the initial state.
 - **Actions:** Left, Right, Suck (Up/Down for larger environments).
 - **Transition Model:** Defines effects of actions.
 - **Goal Test:** All squares must be clean.
 - **Path Cost:** Each step costs 1.
-

Real-World Problems

Route Finding

- Defined by **locations** and **transitions** between them.
- Applications:
 - Routing in computer networks
 - Automated travel advisory systems
 - Airline travel planning

Traveling Salesperson Problem (TSP)

- **Goal:** Find the shortest tour that visits all cities.

- **Challenges:**
 - Requires keeping track of visited cities.
 - **NP-hard** problem; extensive research exists to optimize solutions.
-

Example: Romania

- You are on holiday in Romania, currently in Arad.
- Your **goal**: Reach Bucharest before your flight tomorrow.

Problem Formulation

- **States**: Various cities.
- **Actions**: Drive between cities.
- **Solution**: A sequence of cities, e.g., *Arad* → *Sibiu* → *Fagaras* → *Bucharest*.

What is a Solution?

- A **sequence of actions** that transforms the initial state into a goal state.
 - Sometimes, the solution is just the goal state itself (e.g., molecular structure inference from mass spectrometry).
-

State Space Representation

- The real world is **complex** → We must abstract the state space for problem-solving.
 - **Abstract State** = Set of real states.
 - **Abstract Action** = Complex combination of real actions.
 - **Abstract Solution** = A set of real paths that are solutions in the real world.
-

Search Concepts

Expanded Nodes vs. Frontier

- **Frontier:** Set of nodes that have been reached but not yet expanded.
- **Interior:** Nodes that have been expanded.
- **Exterior:** States that have not been reached.

Example: The 8-Puzzle

- **States:** Locations of tiles.
 - **Actions:** Move left, right, up, down.
 - **Goal Test:** Check if the state matches the goal state.
 - **Path Cost:** Each move costs 1.
-

Search Elements & Structure

Search Tree Elements

- **Root Node:** Starting state of the problem.
- **Branches (Edges):** Possible actions leading to new states.
- **Nodes:** Represent states in the search space.
- **Leaf Nodes:** Terminal states with no further expansion.
- **Goal Node:** A leaf node that satisfies the goal condition.
- **Path Cost:** Cumulative cost of reaching a node from the root.

Search Algorithm Basics

- The **search tree** is built from the initial state, with actions forming branches.
 - The **search strategy** determines the order in which states are expanded.
-

Search Algorithm Infrastructure

- Each node (n) contains:

- **n.STATE:** The state in the search space.
- **n.PARENT:** The parent node in the search tree.
- **n.ACTION:** The action applied to generate the node.
- **n.PATH-COST:** The cost of reaching this node.

Implementation: States vs. Nodes

- A **state** represents a physical configuration.
 - A **node** is part of the search tree, containing state, parent node, action, path cost, and depth.
 - The **Expand function** generates new nodes based on the given state.
-

Search Strategies

Defining a Search Strategy

- A search strategy determines the order of node expansion.
- **Evaluation Criteria:**
 1. **Completeness** – Does it always find a solution if one exists?
 2. **Time Complexity** – Number of nodes generated.
 3. **Space Complexity** – Maximum number of nodes in memory.
 4. **Optimality** – Does it always find the least-cost solution?

Complexity Notation

- **b** = Maximum branching factor of the search tree.
 - **d** = Depth of the least-cost solution.
 - **m** = Maximum depth of the state space (possibly infinite).
-

This refined format keeps all original content intact while improving readability for exam preparation. Let me know if you need any additional edits! 🚀

Problem Solving by Searching

Contents

- Problem Solving by Searching
 - Problem Formulation
 - Search Strategies
 - Uninformed Search Strategies
-

Uninformed Search Strategies

Uninformed search strategies explore the search space without prior knowledge about the goal's location, only using the problem definition.

Types of Uninformed Search:

1. Breadth-First Search (BFS)
 2. Uniform Cost Search (UCS)
 3. Depth-First Search (DFS)
 4. Depth-Limited Depth-First Search (DLS)
 5. Iterative Deepening Depth-First Search (IDDFS)
-

Breadth-First Search (BFS)

- Explores all nodes at the current depth before moving deeper.
- Uses a **queue (FIFO)** for node expansion.

- **Guaranteed** to find the shortest path in an unweighted graph.
- **High memory consumption** as it stores all nodes at a given depth.

Properties of BFS:

- **Branching factor (b):** Number of successors per node.
- **Level progression:** Root \rightarrow b nodes \rightarrow b^2 nodes \rightarrow b^3 nodes \rightarrow ... \rightarrow b^d nodes
- **Total nodes at depth d:**

$$1 + b + b^2 + b^3 + \dots + b^d$$

- **Time and Space Complexity:** $O(b^d)$
- **Completeness:** Yes (if b is finite)
- **Optimality:** Yes (if cost = 1 per step)
- **Space Complexity:** The major issue (higher than time complexity).

Real-World Example Breakdown

- **Branching factor (b) = 10**
- **Processing speed:** 1 million nodes/sec
- **Memory per node:** 1 KB
- **At depth d = 10:**
 - Time: < 3 hours
 - Memory required: **10 TB**
- **At depth d = 14:**
 - Processing time \approx **3.5 years**

Challenges:

- **Exponential complexity** makes BFS impractical for:
 - Large datasets
 - Deep searches
 - High branching factors

Solution:

Use **informed search strategies** for real-world applications.

Uniform Cost Search (UCS)

- Also called **Dijkstra's Algorithm** in theoretical computer science.
 - Unlike BFS, which expands nodes in increasing depth order, UCS expands nodes in order of **path cost**.
-

Depth-First Search (DFS)

- Explores as deep as possible along a branch before backtracking.
- Uses a **stack (LIFO)** for node expansion.

Properties of DFS:

- **Completeness:** No (fails in infinite-depth spaces or loops).
 - Solution: Modify to avoid repeated states → complete in finite spaces.
- **Optimality:** No (not guaranteed to find the shortest path).
- **Time Complexity:** $O(b^m)$, where m = max depth.
- **Space Complexity:** $O(bm)$

Backtracking Optimization:

- Instead of generating all successors at once, **backtracking** generates **one successor at a time**.
 - Reduces memory usage to **$O(m)$** instead of **$O(bm)$** .
-

Depth-Limited Depth-First Search (DLS)

- DFS with a **predefined depth limit (L)**.
- Prevents infinite loops by stopping at the depth limit.
- If the goal is beyond L, it **may fail**.

Complexity:

- Time Complexity: $O(b^L)$
- Space Complexity: $O(bL)$

Example Use Case:

- On a map of Romania, there are **20 cities**.
 - If a solution exists, the longest path must be **at most 19 steps**.
 - Setting $L = 19$ ensures completeness.
-

Iterative Deepening Depth-First Search (IDDFS)

- Repeatedly applies DLS with increasing depth limits.
- Combines DFS's low memory use with BFS's completeness.
- Guaranteed to find the shortest path.

Complexity:

- Time Complexity: $O(b^d)$ (some nodes are re-explored but remains efficient).
- Space Complexity: $O(bd)$

When is IDDFS Used?

- Commonly applied in **large search spaces** where BFS's memory usage is impractical.

IDDFS Algorithm:

- Repeatedly applies **depth-limited search** with **increasing limits**.
- Terminates when:
 1. A solution is found.
 2. The depth-limited search returns failure (no solution exists).

Node Generation Comparison:

- Depth-Limited Search (DLS) to depth d with branching factor b :

$$N_{DLS} = 1 + b + b^2 + \dots + b^{d-1} + b^d$$

- Iterative Deepening Search (IDS) to depth d :

$$N_{IDS} = (d + 1)b^0 + db^1 + (d - 1)b^2 + \dots + 3b^{d-2} + 2b^{d-1} + 1b^d$$

- For $b = 10$, $d = 5$:
 - DLS nodes generated: 111,111
 - IDS nodes generated: 123,456
 - Overhead: 11%

Properties of IDDFS:

- Completeness: Yes
 - Time Complexity: $O(b^d)$
 - Space Complexity: $O(bd)$
 - Optimality: Yes (if step cost = 1)
-

Bidirectional Search

- Searches **simultaneously from the start and from the goal**, hoping the two searches meet.
- **Significant reduction in search complexity:**
 - Instead of $O(b^d)$, it requires $O(b^{(d/2)} + b^{(d/2)})$, which is significantly smaller.

Example:

- Find a path from vertex **0** to **14**.
- Two searches are executed:
 - One from vertex **0**
 - One from vertex **14**
- When they meet at **vertex 7**, the search terminates.

Bidirectional Best-First Search

- Maintains **two frontiers** and **two tables of reached states**.

- The node expanded **next** is the one with the **minimum evaluation function** from either frontier.
 - **Once paths meet, they are joined to form a solution.**
 - **The first solution found is not necessarily the best**—search continues until termination criteria are met.
-

Summary

- **Problem formulation** involves abstracting real-world details into a **state space**.
 - **Uninformed search strategies** explore the search space without **prior knowledge** of the goal.
 - **Iterative deepening search** is often the best uninformed search method:
 - Uses **only linear space**
 - Requires **not much more time** than other uninformed algorithms.
-

This version keeps all the content intact but makes it **well-structured, readable, and easy to review** for your exam preparation. Let me know if you need further refinements! 🚀