

# **Message Passing Interface (MPI)**

## **Basics**

Usama Bin Amir

Department of Cyber Security,  
Air University, Islamabad

# Distributed Systems - Definition

---

A **distributed system** is a **software system** in which **components** located on **networked computers** **communicate** and **coordinate** their actions **by passing messages** (to achieve a common goal)

## How to Program Distributed Computers?

- **Message Passing based Programming Model**

# MPI (Message Passing Interface)?

---

- **Standardized** message passing library **specification**
  - for **parallel computers clusters**
  - not a specific product, compiler specification etc.
  - many implementations, **MPICH**, **LAM**, **OpenMPI**  
...
- **Portable**, with **Fortran** and **C/C++** interfaces
- **Real** parallel programming

# A Brief History - MPI

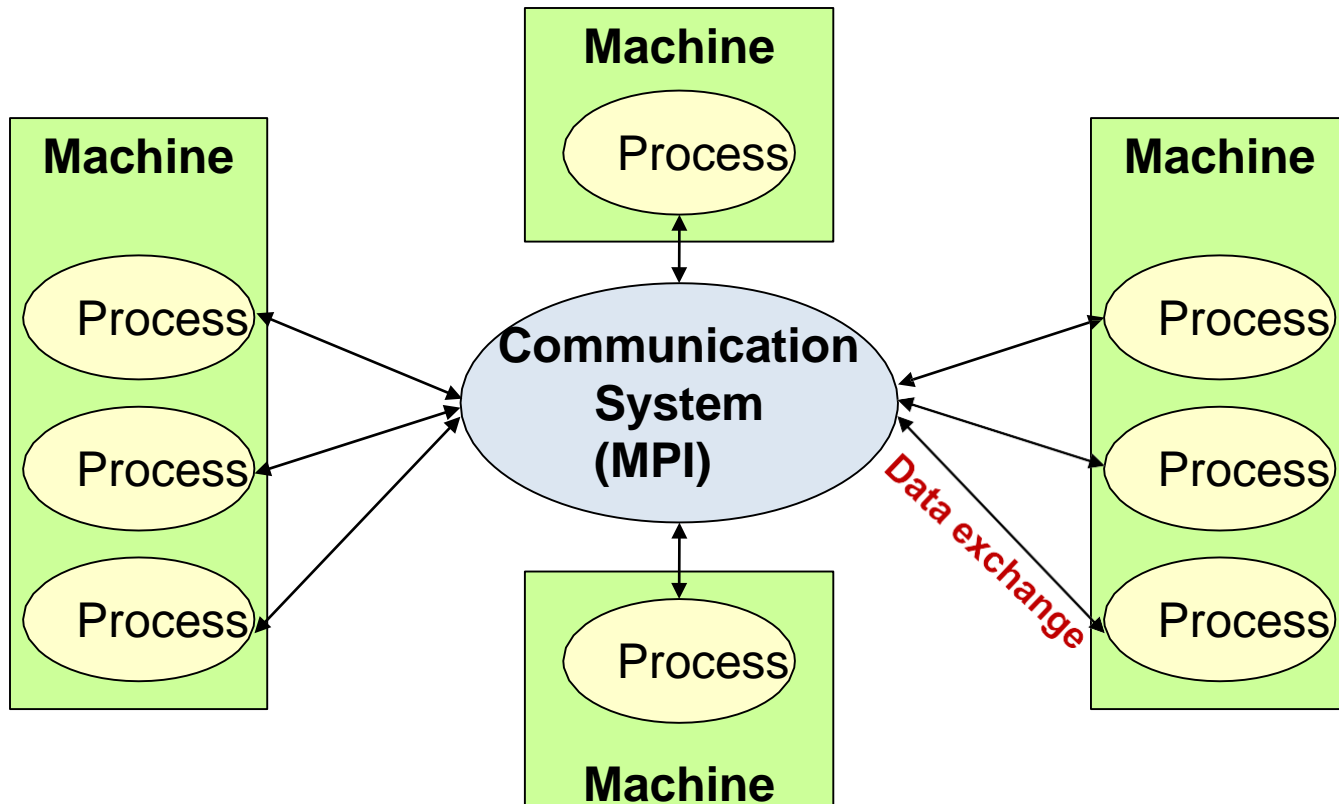
---

- Writing parallel applications is fun!!
- Initially, it was a difficult:
  - No single standard
  - Various implementations (with different features)
- **Solution**: a **MPI** standard was defined
  - Supporting **same features** and **semantics** across implementations
  - By 1994, a **complete interface** and **standard** was defined (**MPI-1**)
- **Result**: **Portable Parallel Programs**

# The Message-Passing Model

## Two major requirements:

1. Creating **separate processes** for execution on different computers
2. Method of **sending** and **receiving** messages



# The Message-Passing Model

---

- A **process** is (traditionally) a **program counter** and **address space**
- **Processes** may have **multiple threads**
  - program counters and associated stacks
  - sharing a single address space
- **MPI** is for **communication** among processes
  - separate address spaces
- **Inter-process communication** consists of:
  - Synchronization
  - Data movement from one process's address space to another's

# Types of Parallel Computing Models

---

## 1. Data Parallel

—Same instructions are carried out simultaneously on multiple data items (e.g., SIMD)

## 2. Task Parallel

—Different instructions on different data (e.g., MIMD)


## 3. SPMD (Single Program, Multiple Data)

—Not synchronized at individual instruction level (e.g., MIMD style execution)

→ **MPI** is for **MIMD/SPMD** type of parallelism


# The Message-Passing Programming Paradigm

## Sequential Programming Paradigm:



A processor may  
run many processes


## Message-Passing Programming Paradigm





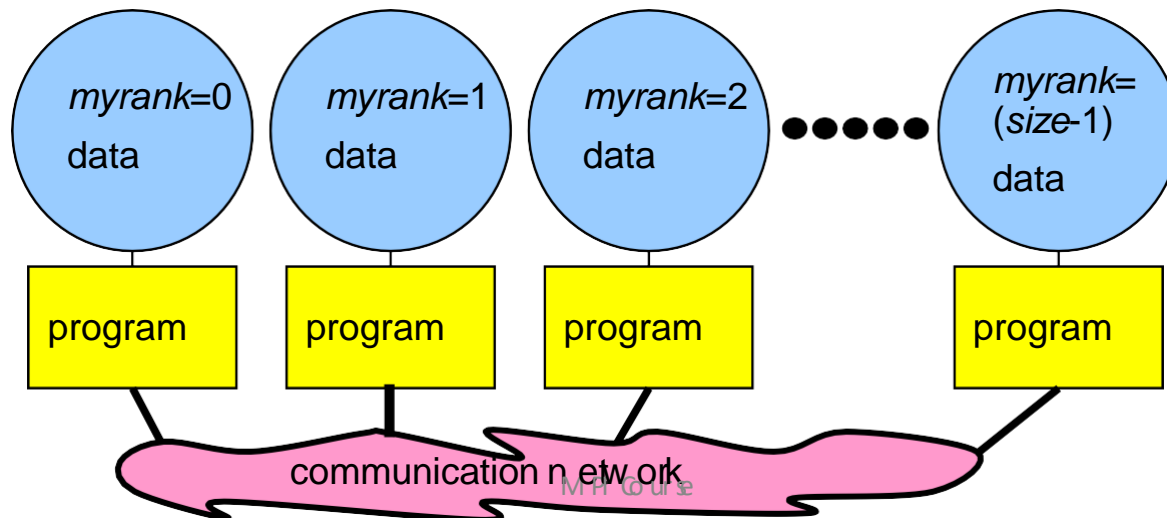
# The Message-Passing Programming Paradigm

- A **process** is a **program** performing a task on a **processor**
- Each **processor/process** runs a **instance/copy** of the **same program**:
  - the **variables** of each sub-program have:
    - the **same name** but **different locations** (distributed memory) and **different data**
    - i.e., all variables are **local to a process**
  - **communicate** via *message passing*




# Data and Work Distribution

- To **communicate** together **MPI-processes** need identifiers: **rank** = identifying number
- Processes are identified using *rank*



# MPI Fundamentals


- A **communicator** defines a **group of processes** that have the **ability** to **communicate** with **one another**
- In a group, each **processes** is **assigned** a **unique rank**
- Use **rank** to **explicitly communicate** with **one another**



# Communicators

**Communicators** do **not** need to contain **all processes** in the system

**Every process** in a communicator has an ID called as “**rank**”




When you **start** an **MPI program**, there is **one predefined communicator** **MPI\_COMM\_WORLD**

The **same process** might have **different ranks** in **different communicators**

**Simple programs** typically **only use** the **predefined communicator** **MPI\_COMM\_WORLD**

# Communicator MPI\_COMM\_WORLD

- All processes of an MPI program are members of the default communicator **MPI\_COMM\_WORLD**
- **MPI\_COMM\_WORLD** is predefined
- Additional communicators can also be defined
- Each process has its own rank in a communicator:
  - starting with 0
  - Ending with (size-1)



# How big is the MPI library?

- **Huge** (125 Functions) !!
- **Good news** – you can write **useful MPI programs** only using **6 basic functions**

---

<code>MPI_Init</code>	Initializes MPI.
<code>MPI_Finalize</code>	Terminates MPI.
<code>MPI_Comm_size</code>	Determines the number of processes.
<code>MPI_Comm_rank</code>	Determines the label of calling process.
<code>MPI_Send</code>	Sends a message.
<code>MPI_Recv</code>	Receives a message.

---

# MPI – Start & Termination

---

**int MPI\_Init(int \*argc, char \*\*\*argv)**

- **initialization:** must **call** this **prior** to **other MPI routines** (by main thread)
- **initializes MPI environment**

**int MPI\_Finalize( )**

- **Must call** at the **end of the computation** (by the main thread)
- **performs various clean-up tasks to terminate MPI environment**
- **Return codes** (for both *MPI\_Init* & *MPI\_Finalize*)
  - **MPI\_SUCCESS**
  - **MPI\_ERROR**

# Communicators

---

## Communicator: MPI\_Comm

- Group of **processes** that could **communicate** with **one another**
- **Communication domains could overlap**
- **A process may be part of multiple communicators**
- **MPI\_COMM\_WORLD:** **root communicator** (all the processes)



# Communicators

---

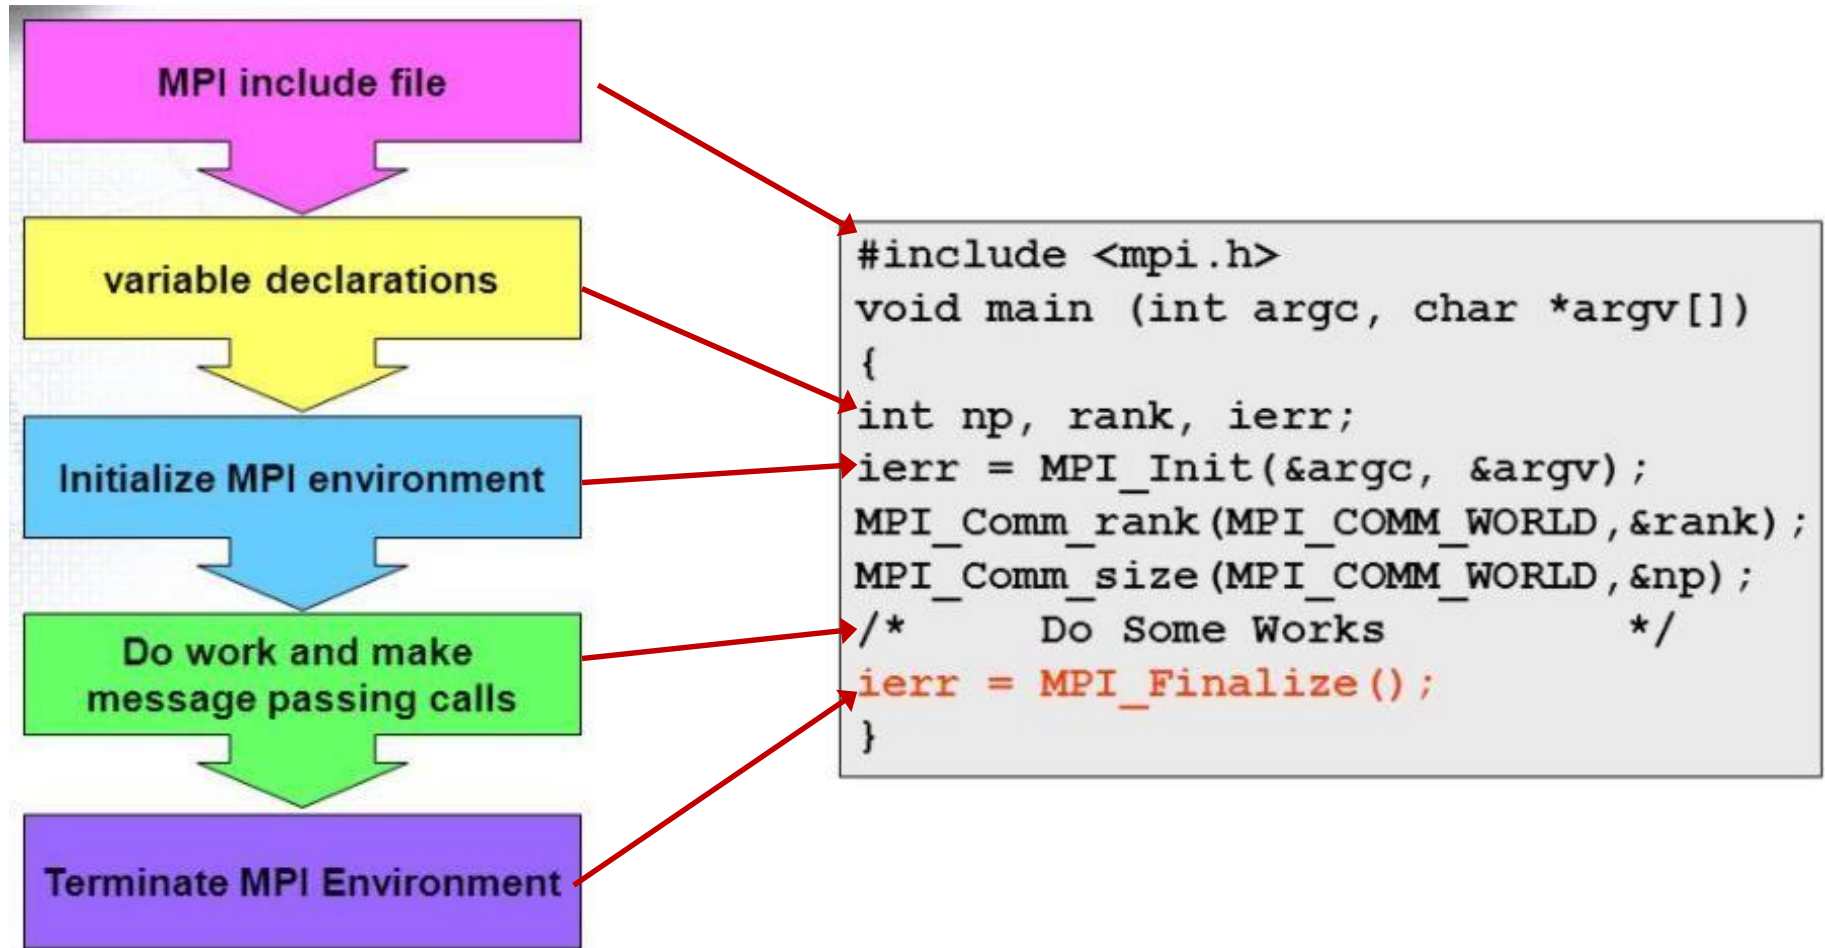
**int MPI\_Comm\_size(MPI\_Comm comm, int \*size)**

- **Determine the number of processes** (in a particular communicator)

**int MPI\_Comm\_rank(MPI\_Comm comm, int \*rank)**

- **Index of the calling process** (in a particular communicator)
- **$0 \leq \text{rank} < \text{communicator size}$**

# MPI Program – A Generic Structure



# A minimal MPI Program

Demo:  
hello.c

```
#include <mpi.h>
#include <stdio.h>
int main(int argc, char *argv[])
{
    MPI_Init(&argc, &argv);
    printf("Hello PDC Class!\n");
    MPI_Finalize();
    return 0;
}
```

MPI header file

```
graph TD; A[MPI header file] --> B[MPI environment initialization...]; B --> C[MPI_Comm_size...]; C --> D[MPI_Comm_rank...]; D --> E[MPI_Get_processor_name...]; E --> F[Demo: processorName.c]; F --> G[clean up the MPI environment];
```

**MPI environment initialization**, all of MPI's **global** and **internal variables** are constructed. For example, a **communicator** is **formed** around all of the **processes** that were **spawned**, and **unique ranks** are **assigned** to each process.

MPI\_Comm\_size returns the **size** of a **communicator**. Here, MPI\_COMM\_WORLD encloses all of the processes, so this call should return the amount of processes that were requested for the job.

MPI\_Comm\_rank returns the rank of a process in a communicator. Ranks are incremental starting from zero and are primarily used for identification purposes during send/receive.

MPI\_Get\_processor\_name obtains the actual name of the processor on which the process is executing.

**Demo:**  
**processorName.c**

clean up the MPI environment

# Compiling MPICH Program

---

- Regular C applications:

```
gcc hello_world.c -o hello_world
```

- MPI based C applications

```
mpicc mpi_hello_world.c -o mpi_hello_world
```

# Running MPICH Program

---

- Regular C applications

```
./hello_world
```

- MPI based C applications (running with 16 processes)

```
mpiexec -n 16 ./mpi_hello_world
```

# Running MPICH Program on a Cluster

---

- **On Clusters**, you will **have to set up a host file** (named `machinefile` in our earlier demo)
- The **host file** contains **names** of all of the **nodes** on which your **MPI job will execute**
- Example (*machinefile* contents):

*slave1:4 # this will spawn 4 processes on slave1*

*master:2 # this will spawn 2 processes on master*

- and executed it using

**`mpiexec -n 6 -f machinefile ./mpi_hello`**

# Running MPICH Program on a Cluster

---

- As shown, the **MPI program** was **launched** across **all** of the **hosts** in *machinefile*
- Each process was **assigned** a **unique rank**, which was printed off **along with** the **process name**.
- The **output** of the **processes** is in an **arbitrary order** since there is **no synchronization involved** before printing



# Congratulations !!

---

- You now have a **basic understanding** about **how MPI works**
- **Even better, you already have a cluster** and you can write parallel programs !!

# **Point-to-Point Communication**

# Data Messages

---

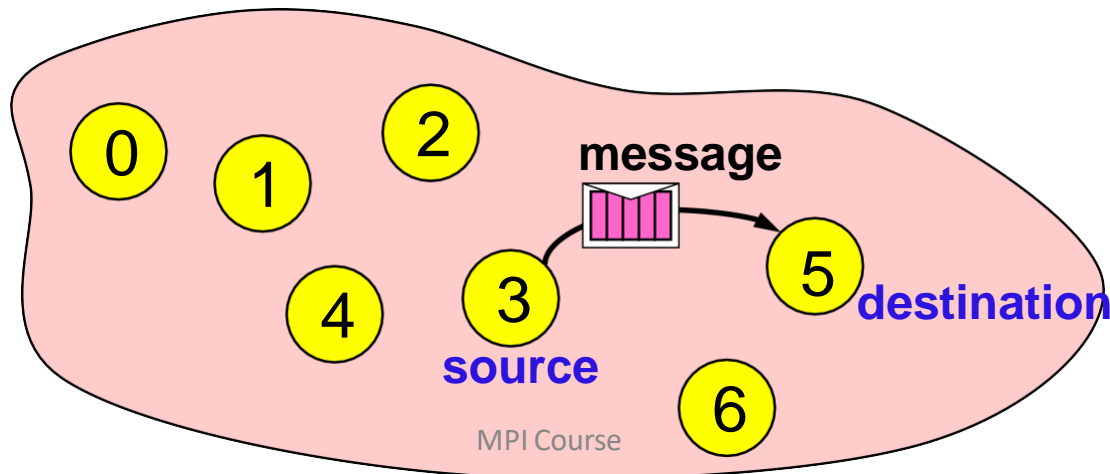
- A **message** contains a **number of elements** of some particular datatype

**Example: message with 5 integers**

2345	654	96574	-12	7676
------	-----	-------	-----	------

# Point-to-Point Communication

- **Communication** between two processes
- **Source** process **sends message** to **destination** process
- **Communication** takes place **within a communicator**, e.g., MPI\_COMM\_WORLD
- **Processes** are **identified** by their **ranks** in the communicator



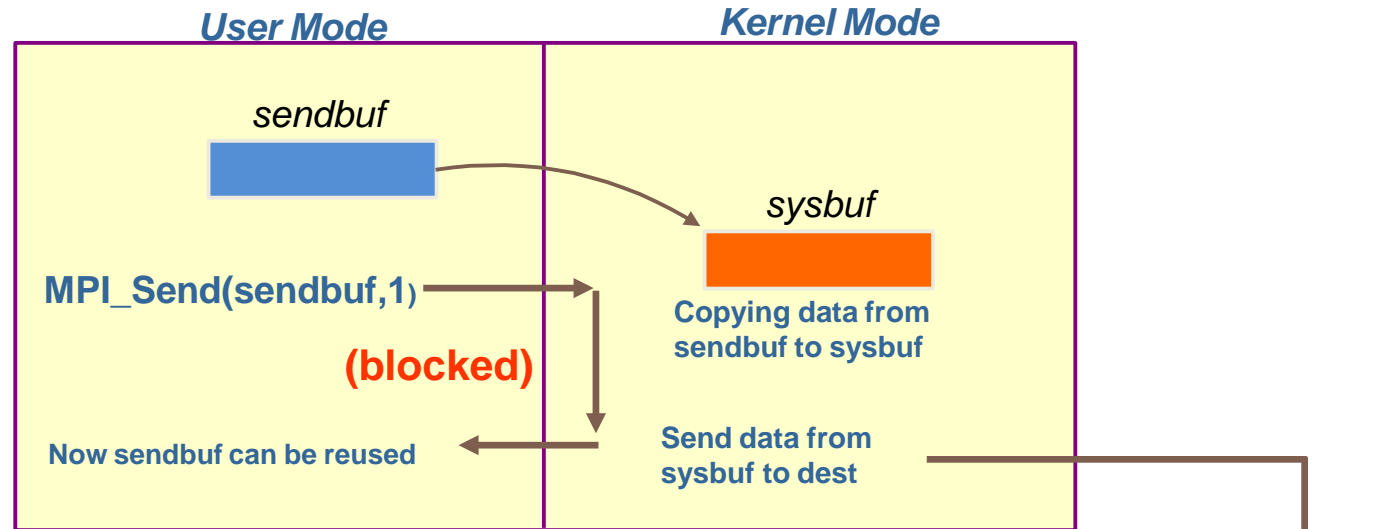
# Point to Point Communication

---

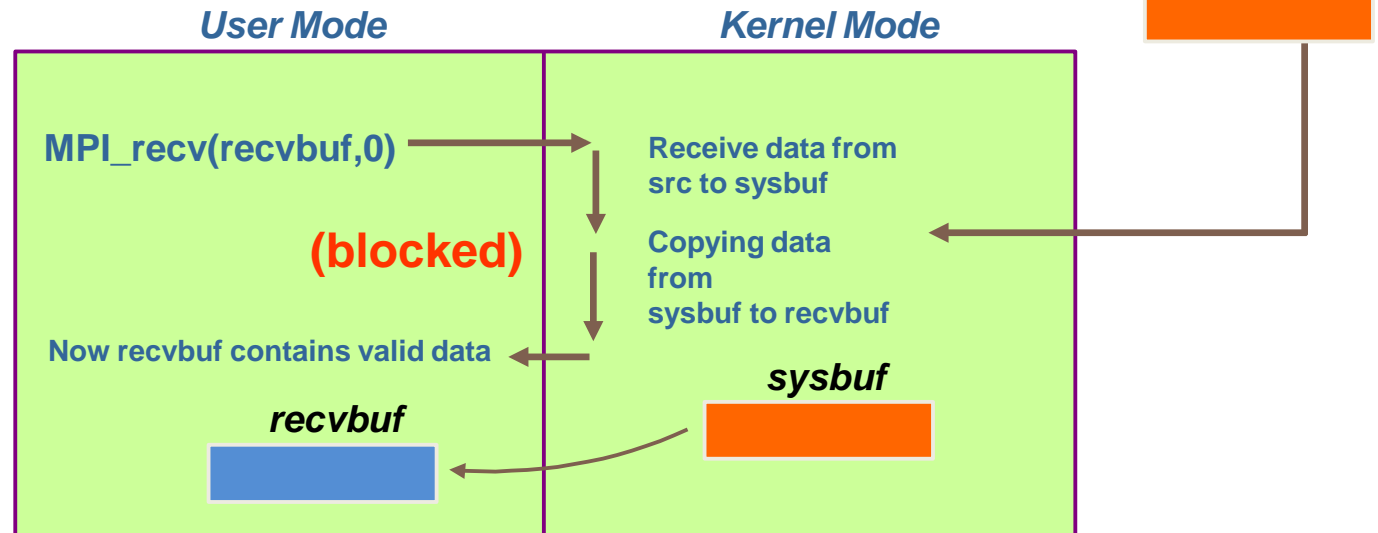
- **Communication** is **done** using **send** and **receive** among processes:
  - To **send** a **message**, **sender** provides the **rank** of the **process** and a **unique tag** to identify the message
  - The **receiver** can then **receive** a **message** with a **given tag** (or it may not even care about the tag), and then **handle the data accordingly**
  - Two basic (and simple) functions, **MPI\_Send** and **MPI\_Recv**

# MPI\_Send & MPI\_Recv

**Sending process** waits until all data are transferred to the system buffer



**Receiving process** waits until all data are transferred from the system buffer to the receive buffer



# Data Communication in MPI

---

- Communication requires the following information:
  - **Sender** has to know:
    - **Whom to send** the data to (*receiver's process rank*)
    - What **kind of data** to send (*100 integers* or *200 characters*,...)
    - A user-defined “**tag**” (**distinguish different messages**)
  - **Receiver** “**might**” have to know:
    - **Who is sending** Or wildcard: **MPI\_ANY\_SOURCE** (meaning anyone can send)
    - **kind of data** is being **received** (may be partial info, e.g., upper bound)
    - Message “**tag**” Or wildcard: **MPI\_ANY\_TAG** (any message)

# MPI Send

---

```
MPI_Send(void* data, int count, MPI_Datatype  
type, int dest, int tag, MPI_Comm comm)
```

***data***: pointer to data

***count***: number of elements to send

***type***: data type of data

***dest***: destination process (rank)

***tag***: identifying tag

***comm***: communicator

- **Blocking operation**
- When **MPI\_Send** returns the **message is sent** and the **data buffer can be reused** (the message may not have been received by the target process yet)



# MPI\_Recv

---

`MPI_Recv`(void\* **data**, int **count**, MPI\_Datatype **type**, int **source**, int **tag**, MPI\_Comm **comm**, MPI\_Status\* **status**)

***data***: pointer to data

***count***: number of elements to be received (upper bound)

***type***: data type

***source***: source process of the message

***tag***: identifying tag

***comm***: communicator

***status***: i.e., *sender*, *tag*, and *message size*

- When `MPI_Recv` returns the message has been received
- **Waits** until a **matching** (on **source**, **tag**, **comm**) **message** is **received**

# Elementary MPI datatypes

---

Similar to C datatypes, portable MPI datatypes

int → MPI\_INT

double → MPI\_DOUBLE

char → MPI\_CHAR

- **Complex datatypes** also possible:
  - E.g., a **structure datatype** that **comprises** of other datatypes → a **char**, an **int** and a **double**

# Elementary MPI datatypes

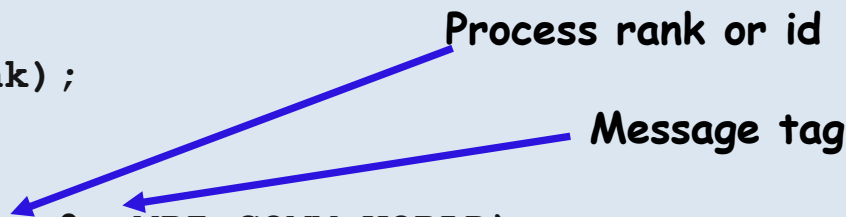
MPI data type	C data type
MPI_CHAR	signed char
MPI_SHORT	signed short int
MPI_INT	signed int
MPI_LONG	signed long int
MPI_UNSIGNED_CHAR	unsigned char
MPI_UNSIGNED_SHORT	unsigned short int
MPI_UNSIGNED	unsigned int
MPI_UNSIGNED_LONG	unsigned long int
MPI_FLOAT	float
MPI_DOUBLE	double
MPI_LONG_DOUBLE	long double
MPI_BYTE	8 bits
MPI_PACKED	packed sequence of bytes

# Simple Communication in MPI

Demo: P2P.c

```
#include <mpi.h>
#include <stdio.h>
int main(int argc, char **argv)
{
    int rank;
    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    if (rank == 0){
        char sdata[] = "Hello PDC";
        MPI_Send(sdata, 9, MPI_CHAR, 1, 0, MPI_COMM_WORLD);
    }
    else if (rank == 1) {
        char rdata[]="";
        MPI_Recv(rdata,9,MPI_CHAR,0,0,MPI_COMM_WORLD,MPI_STATUS_IGNORE);
        printf("\nI am a slave,received %s message from master\n", rdata);
    }

    MPI_Finalize();
    return 0;
}
```



Process rank or id

Message tag

# The MPI\_Status structure

- Information is returned from **MPI\_RECV** in *status*

```
typedef struct MPI_Status {  
    int MPI_SOURCE;  
    int MPI_TAG;  
    int MPI_ERROR;  
};
```

- For the **messages** received using **wildcards** (**MPI\_ANY\_TAG**, **MPI\_ANY\_SOURCE**)
  - Receiver can check **actual** “**Source**” and “**Tag**” of the **message**
  - MPI\_STATUS\_IGNORE** can be used if we don't need any additional **information**

# The MPI\_Status structure

---

```
int MPI_Get_count(MPI_Status* status,  
                  MPI_Datatype type,  
                  int* count)
```

- In **MPI\_Get\_count**, the user passes the **MPI\_Status** structure, the **datatype** of the **message**, and **count** is returned
- ***count*** variable is the **total number** of elements that were **received**

```
const int MAX_NUMBERS = 100;
int numbers[MAX_NUMBERS];
int number_amount;
if (world_rank == 0) {
    // Pick a random amount of integers to send to process one
    srand(time(NULL));
    number_amount = (rand() / (float)RAND_MAX) * MAX_NUMBERS;

    // Send the amount of integers to process one
    MPI_Send(numbers, number_amount, MPI_INT, 1, 0, MPI_COMM_WORLD);
    printf("0 sent %d numbers to 1\n", number_amount);
} else if (world_rank == 1) {
    MPI_Status status;
    // Receive at most MAX_NUMBERS from process zero
    MPI_Recv(numbers, MAX_NUMBERS, MPI_INT, 0, 0, MPI_COMM_WORLD,
              &status);

    // After receiving the message, check the status to determine
    // how many numbers were actually received
    MPI_Get_count(&status, MPI_INT, &number_amount);

    // Print off the amount of numbers, and also print additional
    // information in the status object
    printf("1 received %d numbers from 0. Message source = %d, "
           "tag = %d\n",
           number_amount, status.MPI_SOURCE, status.MPI_TAG);
}
```

# Summary (Blocking Send/Recv)

**MPI\_SEND** does not return until **buffer is empty** (available for reuse)

**MPI\_RECV** does not return until **buffer is full** (available for use)

**MPI\_Send** is a **blocking operation**. It may not complete until a matching receive is posted.

- **Improper use** may lead to **Deadlocks** too:

*If (rank == 0) Then*

*Call mpi\_send(..)*

*Call mpi\_recv(..)*

Usually deadlocks → *Else*

*Call mpi\_send(..)*      ← UNLESS you reverse send/recv

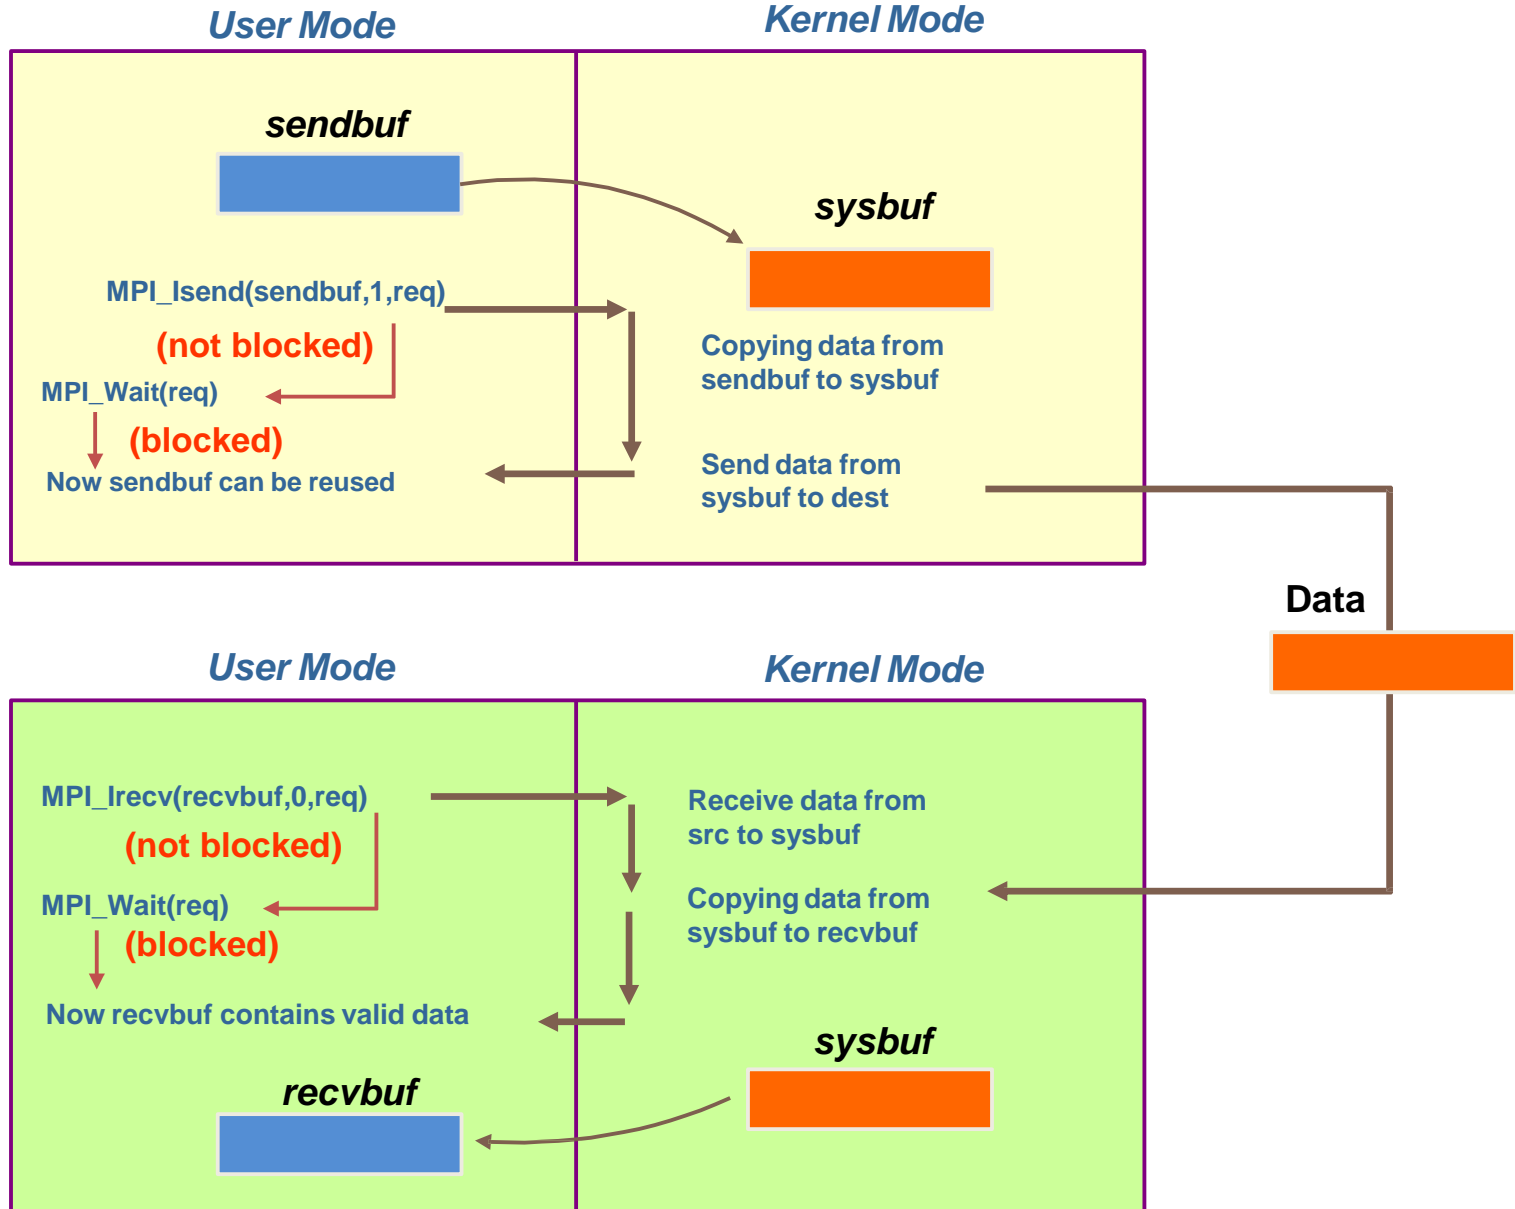
*Call mpi\_recv(..)*

*Endif*




# **Non-Blocking Point-to-Point** **Communication**

# Non-Blocking Communication - Overview




# Non-Blocking Communication - Overview

## Non-blocking Send



## Non-blocking Receive



 = waiting until operation locally completed

# Non-Blocking Send and Receive

---

`MPI_ISEND(buf, count, datatype, dest, tag, comm, request)`

`MPI_Irecv(buf, count, datatype, dest, tag, comm, request)`

*request* is a **request handle** which can be used to query:

- **Check** the status of the communication
- Or **wait** for the completion

# Non-blocking Communication

- **MPI\_Isend** or **MPI\_Irecv** starts communication and returns *request* data structure
- **MPI\_Wait** (also **MPI\_Waitall**, **MPI\_Waitany**) uses *request* as an argument and blocks until communication is complete
- **MPI\_Test** uses *request* as an argument and checks for completion (non-blocking)
- Advantages:
  - No deadlocks (using MPI\_Test for completion check)
  - Overlap communication with computation
  - Exploit bi-directional communication

# Non-Blocking Send and Receive (Cont.)

`MPI_WAIT (request, status)`

`MPI_TEST (request, flag, status)`

Demo:  
P2PNonBlock.c

- `MPI_WAIT` will block until the non-blocking send/receive with the desired request is done
- The `MPI_TEST` is simply queried to see if the communication has completed (`TRUE` or `FALSE`) is returned immediately in `flag`

# Non-blocking Message Passing - Example

```
#include "mpi.h"
#include <stdio.h>

int main(int argc, char* argv[])
{
    int rank, size;
    int tag, destination, count;
    int buffer;

    tag = 1234;
    destination = 1;
    count = 1;
    MPI_Status status;
    MPI_Request request = MPI_REQUEST_NULL;

    MPI_Init(&argc, &argv);

    MPI_Comm_size(MPI_COMM_WORLD, &size);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
```

Demo:  
P2PNonBlock.c

# Non-blocking Message Passing - Example

```
if (rank == 0) { /* master process */
    buffer = 9999;
    MPI_Isend(&buffer, count, MPI_INT, destination, tag, MPI_COMM_WORLD, &request);
}

if (rank == destination) /* slave process */
    MPI_Irecv(&buffer, count, MPI_INT, 0, tag, MPI_COMM_WORLD, &request);

MPI_Wait(&request, &status); //Everyone wait here (both sender & receiver)

if (rank == 0)
    printf("process %d sent %d\n", rank, buffer);

if (rank == destination)
    printf("process %d rcv %d\n", rank, buffer)

MPI_Finalize();
return 0;
}
```



# Non-blocking Message Passing - Example



## Standard Output

```
Compiling  
Compilation is OK  
Execution ...  
processor 0 sent 9999  
processor 1 rcv 9999  
Done.
```

# MPI\_Probe

---

- Instead of posting a receive and simply providing a really large buffer to handle all possible sizes of messages
- You can use **MPI\_Probe** to query the message size before actually receiving it:

```
MPI_Probe(  
    int source,  
    int tag,  
    MPI_Comm comm,  
    MPI_Status* status)
```

# An Example

---

// Probe for an incoming message from process 0, tag 0

```
MPI_Probe(0, 0, MPI_COMM_WORLD, &status);
```

// When probe returns, the status object has the size and other  
// attributes of the incoming message. Get the message size

```
MPI_Get_count(&status, MPI_INT, &number_amount);
```

// Allocate a buffer to hold the incoming numbers

```
int* number_buf = (int*)malloc(sizeof(int) * number_amount);
```

// Now receive the message with the allocated buffer

```
MPI_Recv(number_buf, number_amount, MPI_INT, 0, 0,  
MPI_COMM_WORLD, MPI_STATUS_IGNORE);
```

Demo:  
messageProbe.c

**Any Questions**